




STEVENS
INSTITUTE *of* TECHNOLOGY
THE INNOVATION UNIVERSITY®


CPE/EE 695: Applied Machine Learning

Lecture 2-1: Linear Regression

Dr. Shucheng Yu, Associate Professor
Department of Electrical and Computer Engineering
Stevens Institute of Technology




Polynomial Curve Fitting



$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

Polynomial Regression


Sum-of-Squares Error Function



$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

Minimize $E(\mathbf{w})$ for unknown \mathbf{w} . (maximum likelihood)


Linear Regression



Observation: Data points seem around a line

Let's assume the line is: $\hat{y} = w_0 + w_1 x_1$

Linear Regression



Observation: Data points seem around a line

Let's assume the line is: $\hat{y} = w_0 + w_1 x_1$

Learning the linear model is to estimate the unknown coefficients w_0 and w_1 that make the line fit the data points.



Linear Regression

More generally, we can extend the linear model to (n+1)-dimension:

$$\hat{y} = w_0 + w_1x_1 + w_2x_2 + \cdots + w_nx_n$$

Due to errors/noises (e.g., caused by sensors), observed or measured value of \hat{y} would be:

$$y = w_0 + w_1x_1 + w_2x_2 + \cdots + w_nx_n + \varepsilon$$

Learning the linear model is to estimate the unknown coefficients that **make the linear model fit the noisy data samples**.

For simplicity, we usually use vectorized representation: $\hat{y} = h_w(x) = \mathbf{w}^T \cdot \mathbf{x}$



Linear Regression

Linear model to be learned:

$$\hat{y} = h_w(x) = w^T \cdot x$$

Observed noisy data (x, y) , where:

$$y = w^T \cdot x + \varepsilon$$

Linear Regression

Linear model to be learned:

$$\hat{y} = h_w(x) = \mathbf{w}^T \cdot \mathbf{x}$$

Observed noisy data (x, y) , where:

$$y = \mathbf{w}^T \cdot \mathbf{x} + \varepsilon$$

Performance measure: **mean Square Error (MSE) cost function:**

$$MSE(X, h_w) = \frac{1}{m} \sum_{i=1}^m (\mathbf{w}^T \cdot \mathbf{x}^{(i)} - y^{(i)})^2$$

More frequently, root mean square error (RMSE) = $\sqrt{MSE(W)}$ is used.

Both leads to the same solution:

$$\hat{w} = \operatorname{argmin}(MSE(w)), \text{ i.e., the } w \text{ that minimizes } MSE(w).$$



Linear Regression

The Normal Equation

Actually, there is a closed-form solution (the normal equation) to the problem.

$$\hat{\mathbf{w}} = (\mathbf{X}^T \cdot \mathbf{X})^{-1} \cdot \mathbf{X}^T \cdot \mathbf{y}$$


where $\hat{\mathbf{w}}$: the value of that minimizes the cost function

\mathbf{X} : the training data set, i.e., the matrix $(x^{(1)} \ x^{(2)} \dots \ x^{(m)})^T$.

\mathbf{y} : the vector of target values containing $(y^{(1)} \ y^{(2)} \dots \ y^{(m)})^T$.


The complexity for computing the normal equation is about $O(n^{2.4})$ to $O(n^3)$ where n is the number of features. When n is large (e.g., tens of thousands), computing the normal equation will become very slow.

Gradient Descent



$$w^{(next\ step)} = w - \lambda \cdot \nabla_w MSE(w)$$

Gradient Descent



1. Start with random values of w , obtaining a random polynomial:
 $\hat{y} = h_w(x) = \mathbf{w}^T \cdot \mathbf{x}$
2. Compute MSE using \hat{y} and data (x, y)
3. If MSE is not small enough, update w with the gradient descent approach.
4. Repeat Steps 2 and 3.

$$MSE(X, h_w) = \frac{1}{m} \sum_{i=1}^m (\mathbf{w}^T \cdot \mathbf{x}^{(i)} - y^{(i)})^2$$

$$w^{(next\ step)} = w - \lambda \cdot \nabla_w MSE(w)$$



What is Grade Descent?

A first-order iterative optimization algorithm to find the minimum of a multivariable function $F(\mathbf{x})$.

Rational:

If $F(\mathbf{x})$ is differentiable around a point A , F decreases **fastest** from A in the direction of negative gradient of $F(\mathbf{x})$ at A (i.e., $-\nabla F(A)$). In other words, let

$$A_{n+1} = A_n - \lambda * \nabla F(A)$$


for **small enough** λ , we have $F(A_{n+1}) \leq F(A_n)$.

Through a set of such points A_0, A_1, \dots , it converges to a **local minimum**.

If function $F(\mathbf{x})$ is **convex**, the local minimum is the **global minimum**.

Learning Rate

Gradient Descent



A big learning rate may cause the algorithm to diverge without finding a good solution.

A small learning rate may cause the algorithm to converge very slowly.

Appropriate learning rate λ is very important



Gradient Descent

Gradient descent approach computes gradient $\nabla_{w_j} MSE(w) = \frac{\partial}{\partial w_j}$ for w_j at each iteration:

$$\frac{\partial}{\partial w_j} = \frac{2}{m} \sum_{i=1}^m (w^T \cdot x^{(i)} - y^{(i)}) x_j^{(i)}$$

1) Batch Gradient Descent

Using ALL data sets to calculate the gradient and update w , i.e., $m = N$, where N is the number of all data sets.

2) Stochastic Gradient Descent


Using ONE RANDOM sample to calculate the gradient and update w , i.e., $m = 1$.

3) Mini-batch Gradient Descent

Using a small random set of samples, i.e., $m = r$ where $r \ll m$.

Gradient Descent

Three Modes – Batch, Stochastic and Mini-batch GD



Picture generated with Github code of textbook 2 (Aurelien Geron)

In practice, **Mini-batch** GD is widely used because of the good tradeoffs. The size of the “mini batches” is a hyperparameter.

Convergence rate:


- It takes least iterations for the Batch GD to converge
- It takes the most iterations for the Stochastic GD to converge.
- Mini-batch is in the middle.

Computation at each iteration:

- Batch GD uses all training example to compute the gradient
- Stochastic GD uses one random training example to compute the gradient
- Mini-batch is in the middle

Gradient Descent

Gradient Descent:



Local minimum issue


Batch Gradient Descent: more likely to have local minimum
Stochastic Gradient Descent: less likely to have local minimum
Mini-Batch Gradient Descent: in the middle of the two.

Regularized Linear Models

Ridge Regression

Cost function: $E(w) = \sum_{i=1}^m (\mathbf{w}^T \cdot \mathbf{x}^{(i)} - y^{(i)})^2 + \lambda \sum_{i=1}^m w_i^2$

It has a closed-form solution $\mathbf{w} = (\lambda \mathbf{I} + \Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$.




Lasso Regression

Cost function: $E(w) = \sum_{i=1}^m (\mathbf{w}^T \cdot \mathbf{x}^{(i)} - y^{(i)})^2 + \lambda \sum_{i=1}^m |w_i|$

No closed-form solution for \mathbf{w} ;

But it tends to eliminate the weights of least important features.



Elastic Net

Cost function: $E(w) = \sum_{i=1}^m (\mathbf{w}^T \cdot \mathbf{x}^{(i)} - y^{(i)})^2 + \lambda_1 \sum_{i=1}^m w_i^2 + \lambda_2 \sum_{i=1}^m |w_i|$

In the middle of Ridge and Lasso.

Usually is preferred over Lasso or Ridge.



Cost Functions

Mean Square Error (MSE):

$$MSE(X, h) = \frac{1}{m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2$$

Root Mean Square Error (RMSE):

$$RMSE(X, h) = \sqrt{MSE} \quad (\text{Euclidean norm})$$

Mean Absolute Error (MAE):

$$MAE(X, h) = \frac{1}{m} \sum_{i=1}^m |h(x^{(i)}) - y^{(i)}| \quad (\text{Manhattan norm})$$

l_k *norm* of a vector v with n elements: $\|v\|_k = (|v_0|^k + \dots + |v_n|^k)^{\frac{1}{k}}$

The higher the norm index, the more it focuses on large values and neglect small ones. Therefore, RMSE is more sensitive to outliers than MAE.

Basis Function

Convert Non-Linear Model to Linear Model

Many models are non-linear, for example, polynomial function:

$$\hat{y} = w_0 + w_1x + w_2x^2 + \cdots + w_nx^n = \sum_{i=0}^n w_ix^i$$

To utilize linear regression, we can use a so-called *basis function* $\phi_i(x)$ to do a “mapping”:

$$\phi_i(x) = x^i$$

Therefore, the polynomial function is represented as:

$$\hat{y} = \sum_{i=0}^n w_i\phi_i(x)$$

More generally, the basis function $\phi_i(x)$ can be any linear or non-linear function. We can train the model using linear regression. The normal equation becomes:

$$\hat{w} = (\Phi^T \cdot \Phi)^{-1} \cdot \Phi^T \cdot y, \quad \text{where } \Phi = \begin{pmatrix} \phi_0(x_1) & \phi_1(x_1) & \cdots & \phi_m(x_1) \\ \phi_0(x_2) & \phi_1(x_2) & \cdots & \phi_m(x_2) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(x_n) & \phi_1(x_n) & \cdots & \phi_m(x_n) \end{pmatrix}$$



STEVENS
INSTITUTE *of* TECHNOLOGY

THE INNOVATION UNIVERSITY®

stevens.edu