

CS 559 Machine Learning

Lecture 3: Linear Classification

Ping Wang

Department of Computer Science

Stevens Institute of Technology

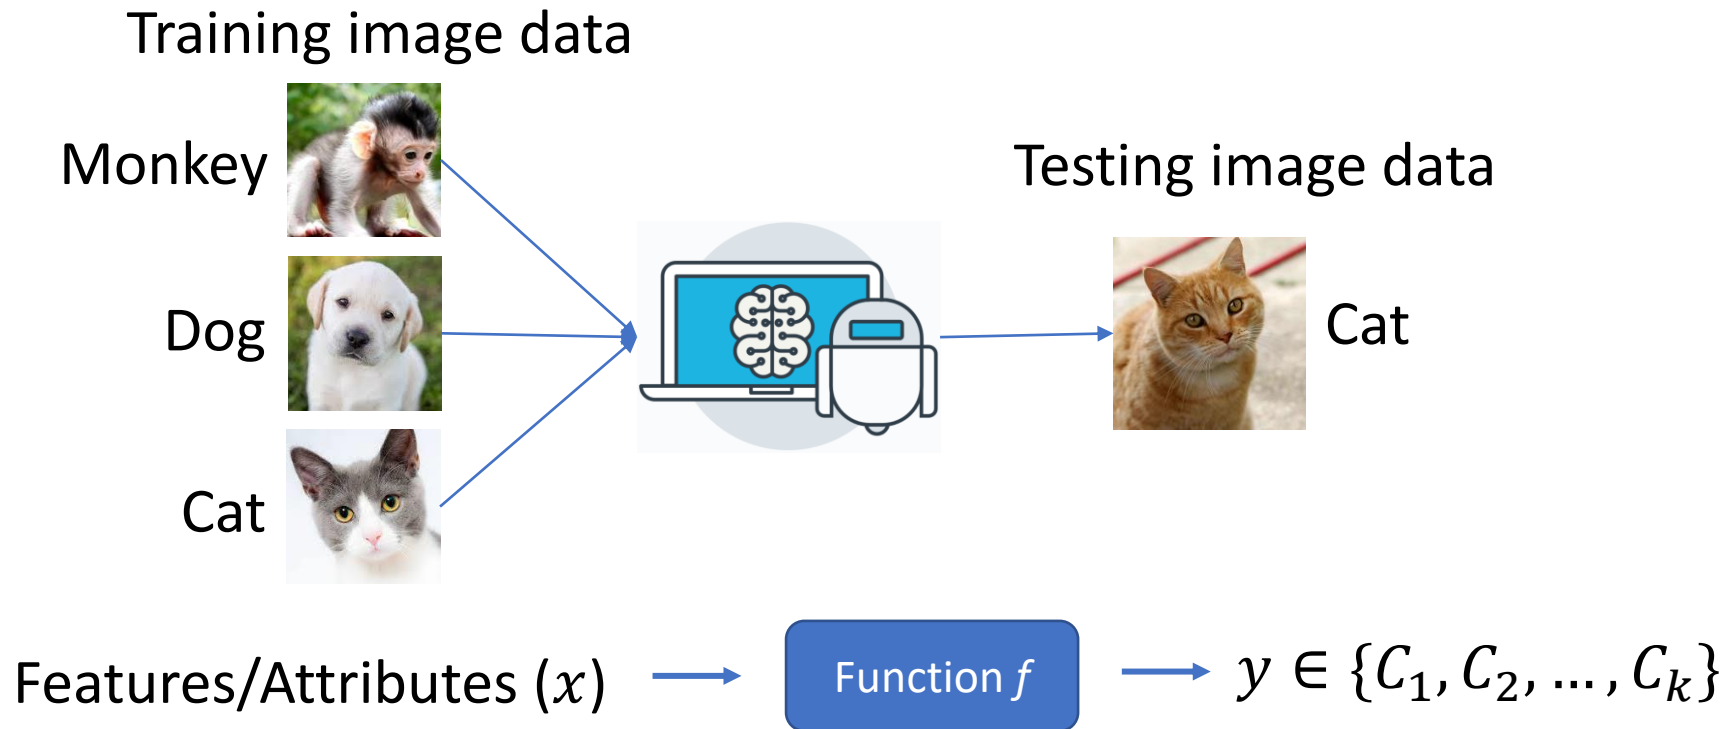


Today's Lecture

- Generative vs Discriminative Classification
- Linear Discriminant Analysis
- The Perceptron Algorithm
- Model Selection

Classification Task

- Task: find a function f that classifies examples into a given set of discrete classes $\{C_1, C_2, \dots, C_k\}$.



Linear Classification

- **Decision boundaries/surfaces**: divide the input space into decision regions.
- For linear classification, the decision boundaries are linear functions of the input x .
- **Linear separable**: datasets that can be exactly separated by linear decision boundaries are linear separable.

Generative and Discriminative Approach

Decision Theory for Classification

- Decision theory, combined with probability theory, allows us to make optimal decisions in situations involving uncertainty.
 - **Training data**: input values X and target values y
 - **Learning stage**: use the training data to learn a model for $p(C_k|x)$, where C_k represents the class k .
 - **Decision stage**: use the learnt posterior probabilities to make optimal class assignments.

Generative Methods

- Solve the inference problem by estimating the **class-conditional density** $p(x|C_k)$ for each class C_k
- Estimate the **class prior probability** $p(C_k)$
- Use Bayes' theorem to get the **class posterior probability**:

$$p(C_k|x) = \frac{p(x|C_k)p(C_k)}{p(x)}$$

$$\text{where } p(x) = \sum_{k=1}^K p(x|C_k)p(C_k)$$

- Use decision theory to determine class label for each new input x .


Discriminative Methods

- **Directly** solve the inference problem of estimating the **class posterior probabilities** $p(C_k|x)$.
- **Discriminative Functions**: Find a function $f(x)$ which maps each input directly onto a class label. Probabilities play no role here.
- Use decision theory to determine class label for each new input x .



Linear Discriminant Function

Binary Classification

- Task: Assign each data point to one of two classes (e.g., 0/1, 1/-1).
- Examples:
 - Is there a face in this image?
 - Is this email spam or not?
 - Based on this brain-scan, does this patient have a given disease or not?
 - Will this customer buy this product or not?
 - Will this patient be re-hospitalized or not?
- Notation: $D = \{(x_1, y_1), \dots, (x_N, y_N)\}$, with $y_n = 1$ if x_n belongs to class 1, and $y_n = -1$ if x_n belongs to class -1.



Linear Discriminant Function



Linear algorithms can be used together with **nonlinear feature spaces** or **nonlinear basis functions** to solve nonlinear classification problems!

Linear Discriminant Function

- Discriminative methods learn the class posterior $p(C_k|x)$.
- The discriminant function here directly assigns the class label for each input vector x .
- The simple way would be the **linear discriminant function**.

Linear Discriminant Function

- Linear discriminants separate the space by a hyperplane, and the parameters define its normal vector.
- Decision function: $f(x) = w^T x + w_0$, where w represents the weight vector, and w_0 is the bias that determines the location of the decision boundary.
- Classification task:
 - If $f(x) \geq 0$, $x \rightarrow$ class 1
 - If $f(x) < 0$, $x \rightarrow$ class -1
- The decision boundary is defined by equation $f(x) = 0$, which is a hyperplane with dimensionality $D - 1$ (D is the dimension of the input space).

Geometrical Properties

- Decision boundary: $f(x) = w^T x + w_0 = 0$
- Let x_1, x_2 be two points which lie on the decision boundary

$$f(x_1) = w^T x_1 + w_0 = 0$$

$$f(x_2) = w^T x_2 + w_0 = 0$$

$$w^T (x_1 - x_2) = 0$$

- Therefore, w represents the orthogonal direction to the decision boundary. It is the normal vector to the hyperplane, and points into the positive class or negative class.

Geometrical Properties: Dot Product

- The decision boundary is the dot product of two vectors w and x , which can be considered as the projection of x onto w .
 - If $f(x) > 0$, $x \rightarrow$ class 1
 - If $f(x) < 0$, $x \rightarrow$ class -1
 - If $f(x) = 0$, x is on the decision boundary.
- Therefore, the decision boundary is the plane that is perpendicular to the weight vector w .

Sign is Important!

- We can observe that the sign of $f(x)$ is important for classification.
- The scale of the weight w does not matter.
- Therefore, we can use the normalized weight w with length $\|w\| = 1$.


Geometrical Properties Cont.

- $w^{*T} = \frac{w^T}{\|w\|}$
- $w^{*T}(x - x_0)$ is the projection of $(x - x_0)$ onto the direction of w^* ; x_0 is a point on the decision boundary.

- Thus

$$\begin{aligned} \frac{w^T}{\|w\|} (x - x_0) &= \frac{1}{\|w\|} (w^T x - w^T x_0) \\ &= \frac{1}{\|w\|} (w^T x + \omega_0) = \frac{f(x)}{\|w\|} \end{aligned}$$

- When $x = 0$, $\frac{f(x)}{\|w\|} = \frac{\omega_0}{\|w\|}$




Signed orthogonal distance of the origin from the decision surface

Linear Discriminant Functions: Multiple Classes

one-versus-the-rest:


- $K - 1$ classifiers, each of which solves a two-class problem of separating points in a particular class C_k from points not in that class.
- Leads to regions of input space that are ambiguously classified, shown in green.



Linear Discriminant Functions: Multiple classes

one-versus-one:

- $\frac{K(K-1)}{2}$ classifiers, one for every possible pair of classes
- Each point is classified by majority voting amongst the discriminant functions.
- Also leads to regions of input space that are ambiguously classified, shown in green.



Linear Discriminant Functions: Multiple classes

- **Solution:** Consider a single K -class discriminant comprising K linear functions of the form

$$f_k(x) = w_k^T x + w_{k0}$$


- Assign a point x to class C_k if

$$f_k(x) > f_j(x) \quad \forall j \neq k.$$

- The decision boundary between class C_k and class C_j is given by:

$$f_k(x) = f_j(x)$$

A hyperplane: $\Rightarrow (w_k - w_j)^T x + (w_{k0} - w_{j0}) = 0$



If two points lie in the same decision region, then any point on the line segment must lie in the same region.

Linear Classification Algorithms

- Mis-classification rate $C(w) = \frac{1}{N} \sum_n \delta[f(x_n) = y_n]$ (i.e. average number of errors) difficult to optimize over w , and might have multiple solutions.
- Many algorithms can be derived by replacing C by another cost-function which can be optimized.
 - Least-square Classification
 - Fisher's Linear Discriminant
 - Rosenblatts' Perceptron
 - Logistic Regression
 - Support Vector Machines


Least Square Classification

Least Square Classification

- We have to fit the function $f(x) = xw^T + \omega_0$ to data.
- Simply do a linear regression from x to y by minimizing the sum-of-squared errors $\sum_n (f(x_n) - y_n)^2$.
- $w_{reg} = (\sum_n x_n x_n^T)^{-1} \sum_n x_n y_n$
- Questions: In what situations might this be a bad idea?

Least Square Classification

- Least squares solutions lack robustness to outliers.
- The decision boundary changed significantly when extra data points are added.




Bishop PRML Figure 4.4.


Purple: least squares. Green: logistic regression

Least Square Classification

- Example of a synthetic data set with three classes.
- The background colors denote the respective classes of the decision regions.



Least square discriminant



Logistic regression

Bishop PRML Figure 4.5

Fisher's Linear Discriminant

Classification via Projection

- A linear function: $f(x) = w^T x + w_0$
assuming in 2D, projects each point x to a line parallel to w :

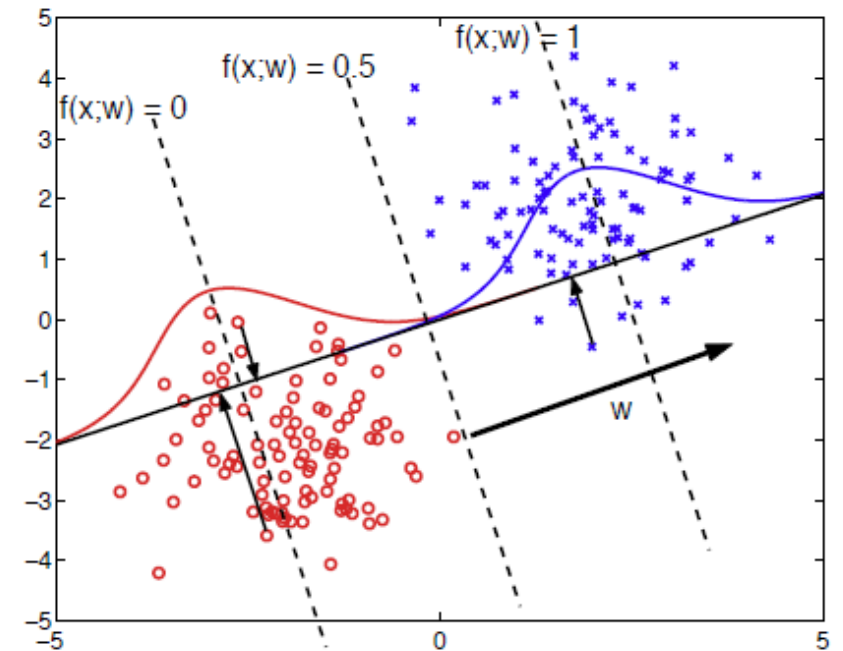
$$x_1 \rightarrow z_1 = w^T x_1$$

$$x_2 \rightarrow z_2 = w^T x_2$$

...





$$x_N \rightarrow z_N = w^T x_N$$

- We can study how well the projected points z_1, \dots, z_N viewed as functions of w are separated across the classes.



Classification via Projection

- By varying w we get different levels of separation between the projected points.
- Find w that maximizes the separation of the projected points across classes.
- Quantify the **separation (overlap)** in terms of means and variances of the resulting 1-dimensional class distributions.



Optimizing the Projection

- Objective:
 - Find w that maximizes the separation of the projected points across classes.
- Solution:
 - Quantify the separation (overlap) in terms of means and variances of the resulting 1-dimensional class distributions.

Fisher's Linear Discriminant

- One way to view a linear classification model is in terms of **dimensionality reduction**.
- For binary classification, suppose we project x onto one dimensional:

$$f = w^T x$$

- A threshold t can be set to assign the label for x :



$$\text{If } f \leq t, x \rightarrow C_1$$

$$\text{If } f > t, x \rightarrow C_2$$

- In general, the projection leads to considerable **loss of information**, and classes well separated in the original space may strongly overlap in one dimension.

Fisher's Linear Discriminant

- Find an orientation along which the projected samples are well separated.
- This is exactly the goal of linear discriminant analysis (LDA).
- In other words: we want to find the linear projection that best separates the data, i.e. best discriminates data of different classes.



Fisher's Linear Discriminant

- We use N_1 and N_2 to represent the number of samples for class C_1 and C_2 , respectively.
- Consider the normalized weight vector w with $\|w\| = 1$.
- Then, $w^T x$ is the projection of x onto the direction of w .
- Objective: find the projections of $w^T x$ for $x \in C_1$ and $x \in C_2$ can be well separated (maximize the class separation).

How to Measure the Separation?

A measure of the separation between the projected points is the **difference of the sample means**.

- Sample mean vector of class C_1 :

$$m_1 = \frac{1}{N_1} \sum_{x \in C_1} x$$



- Sample mean vector of class C_1 after projection:

$$m'_1 = \frac{1}{N_1} \sum_{x \in C_1} w^T x = w^T m_1$$

- Objective is to choose w that can maximize the separation $|m'_1 - m'_2| = w^T |m_1 - m_2|$

How to Measure the Separation?

- Choose w that maximizes projection-distance of class means
 $|m'_1 - m'_2| = w^T |m_1 - m_2|$
- We find that $w \propto m_1 - m_2$
- Limitation: maximizing distance between means, but the projected variances within each class might be large (**large class overlap**).



How to Measure the Separation?

- To obtain good separation of the projected data, we really want the difference between the means to be large relative to some measure of the standard deviation of each class.

- Variance of the projected samples of class C_1 :

$$s_1^2 = \sum_{x \in C_1} (w^T x - m'_1)^2$$

- **Total within-class variance** of the projected samples will be:
 $s_1^2 + s_2^2$

- Fisher linear discriminant analysis: **Fisher criterion** defined by the ratio of the between-class variance to the within-class variance.

$$\arg \max_w \frac{|m'_1 - m'_2|^2}{s_1^2 + s_2^2}$$

Fisher's Linear Discriminant

- Define $J(w) = \frac{|m'_1 - m'_2|^2}{s_1^2 + s_2^2}$. To obtain $J(w)$ as an explicit function of w , we define the following matrices:

$$S_1 = \sum_{x \in C_1} (x - m_1)(x - m_1)^T$$

- Within-class covariance matrix: $S_W = S_1 + S_2$
- Then

$$\begin{aligned} s_1^2 &= \sum_{x \in C_1} (w^T x - m'_1)^2 = \sum_{x \in C_1} (w^T x - w^T m_1)^2 \\ &= \sum_{x \in C_1} w^T (x - m_1)(x - m_1)^T w = w^T S_1 w \end{aligned}$$

Fisher's Linear Discriminant

- Therefore, $s_1^2 = w^T S_1 w$ and $s_2^2 = w^T S_2 w$.

- Further:

$$s_1^2 + s_2^2 = w^T S_1 w + w^T S_2 w = w^T (S_1 + S_2) w = w^T S_W w$$

- Similarly:

$$\begin{aligned} |m'_1 - m'_2|^2 &= (w^T m_1 - w^T m_2)^2 \\ &= w^T (m_1 - m_2)(m_1 - m_2)^T w \\ &= w^T S_B w \end{aligned}$$



- Where $S_B = (m_1 - m_2)(m_1 - m_2)^T$ is the between-class covariance matrix.

Fisher's Linear Discriminant

- Therefore, $J(w) = \frac{|m'_1 - m'_2|^2}{s_1^2 + s_2^2} = \frac{w^T S_B w}{w^T S_W w}$
- Differentiating with respect to w , $J(w)$ is maximized when $(w^T S_B w) S_W w = (w^T S_W w) S_B w$
- We can observe that $S_B w = (m_1 - m_2)(m_1 - m_2)^T w$ always in the direction of $(m_1 - m_2)$ since $(m_1 - m_2)^T w$ is a scalar.
- Also, we only care about the direction of w , so we can drop the scalar factors $(w^T S_B w)$ and $(w^T S_W w)$.
- Therefore, we have the solution:
$$w \propto S_W^{-1}(m_1 - m_2)$$

Summary of Fisher's Linear Discriminant

- $m_1 = \frac{1}{N_1} \sum_{x \in C_1} x, m_2 = \frac{1}{N_2} \sum_{x \in C_2} x$
- Separation: Maximize projection-distance of class means
 $w \propto m_1 - m_2$
- Maximizing distance between means, but the projected variances within each class might be large.
- Fix: Maximize the ratio of between-class variance to within-class variance (“signal to noise”).
- Fisher criterion $J(w) = \frac{|m'_1 - m'_2|^2}{s_1^2 + s_2^2}$
- Updated solution: $w = S_W^{-1}(m_1 - m_2)$



Fisher's Linear Discriminant

- Gives the linear function with the maximum ratio of between-class scatter to within-class scatter.
- The classification problem has been reduced from a d -dimensional problem to a more manageable one-dimensional problem.
- Optimal for multivariate normal class conditional densities.

Fisher's Linear Discriminant: Multi-Class

- The analysis can be extended to multiple classes.
- $S_W = \sum_{k=1}^K \sum_{x_i \in C_k} (x_i - m_k)(x_i - m_k)^T$
- $S_B = \sum_{k=1}^K N_k (m_k - m)(m_k - m)^T$ where m is the global mean; N_k is the number of examples in class k .
- Solve: $S_B v = \lambda S_W v$ the generalized eigenvalue problem
- At most $K - 1$ distinct eigenvalues
- The optimal projection matrix V to a subspace of dimension k is given by the eigenvectors corresponding to the largest k eigenvalues.

Fisher's Linear Discriminant

- LDA is a linear technique for **dimensionality reduction**: it projects the data along directions that can be expressed as linear combination of the input features.
- The “appropriate” transformation depends on the data and on the task we want to perform on the data. **Note that LDA uses class labels.**
- **Non-linear** extensions of LDA exist (e.g., generalized LDA).

The Perceptron Algorithm

The Perceptron Algorithm

- First learning algorithm for neural networks. (Frank Rosenblatt, 1957)
- Originally introduced for character classification, where each character is represented as an image.
- It is a type of linear classifier that makes predictions based on a linear predictor function and a set of weights with the input feature vector.

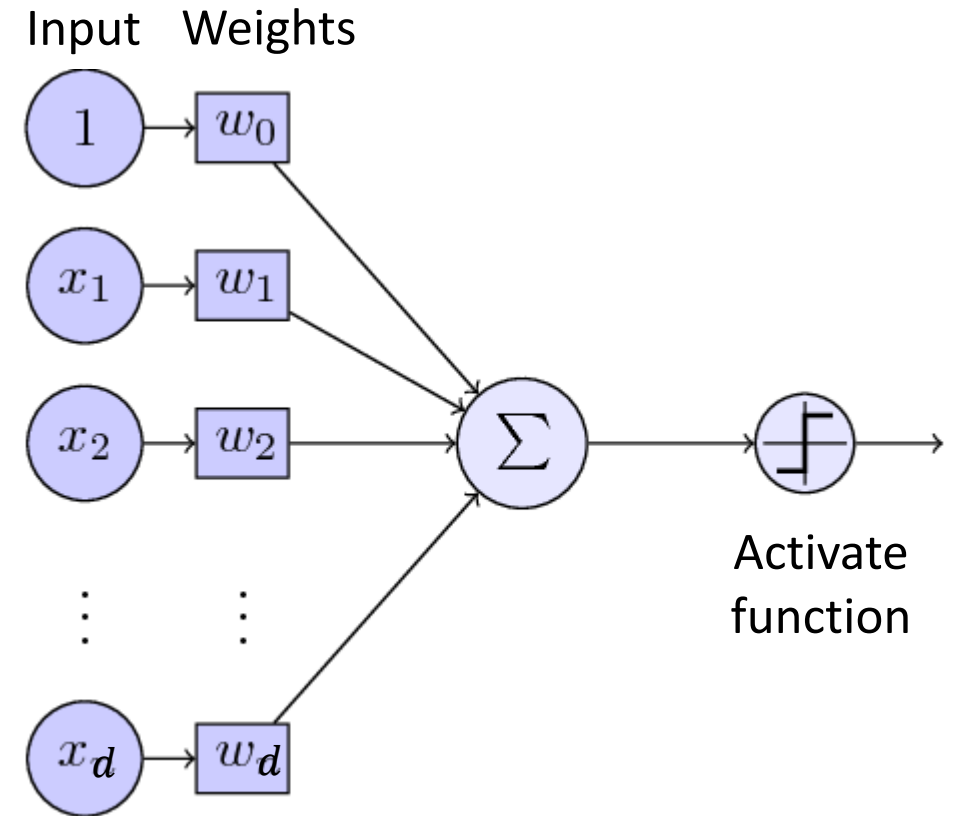
The Perceptron Algorithm

- Total input to output node:

$$\sum_j w_j x_j$$

- Output unit performs the function (activation function):

$$H(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ 0 & \text{if } x < 0 \end{cases}$$



Perceptron: Learning Task

- **Goal:** compute a mapping from inputs to the outputs.
- Example: two-class character recognition problem.
 - Training set: a set of images representing either the character 'a' or the character 'b' (supervised learning);
 - Learning task: learn the weights so that when a new unlabeled image comes in, the network can predict its label.
 - Setting: d input units (intensity level of a pixel), 1 output unit.

Perceptron: Learning Algorithm

- The algorithm proceeds as follows:
 - Initial random setting of weights;
 - The input is a random sequence $\{x_k\}$
 - For each element of class C_1 , if output = 1 (correct), **do nothing**; otherwise, **update weights**.
 - For each element of class C_2 , if output = 0 (correct), **do nothing**; otherwise, **update weights**.

Perceptron: Learning Algorithm

- More formally: $x = (x_1, x_2, \dots, x_d)^T$, $w = (w_1, w_2, \dots, w_d)^T$
- θ : represents the threshold of the output unit
- Unit output: $w^T x = w_1 x_1 + w_2 x_2 + \dots + w_d x_d$
- Output class 1 if $w^T x - \theta \geq 0$
- To eliminate the explicit dependence on θ , output class 1 if $w^T x \geq 0$.

Perceptron: Learning Algorithm

- The objective is to learn the weights so that the perceptron can correctly discriminate elements of C_1 from elements of C_2
- Given x in input, if x is classified correctly, weights are unchanged, otherwise:

$$w = \begin{cases} w + x & \text{if an element of class } C_1 \text{ was classified as in } C_2 \\ w - x & \text{if an element of class } C_2 \text{ was classified as in } C_1 \end{cases}$$

Perceptron: Learning Algorithm

- It is **online**:
 - Only process one example at a time, instead of considering the entire dataset at the same time.
- **Error-Driven Updating**:
 - If it is doing well, it doesn't update the parameters.
 - Only when the prediction is incorrect, it updates the parameters.
 - The parameters are updated in a way that **it would do better on this example next time around**.
- Simple, but works well!

Perceptron: Learning Algorithm

- **1st case:** $x \in C_1$, but was classified in C_2 . In other words, the correct answer is 1, which corresponds to $w^T x \geq 0$, but the model provides $w^T x < 0$. In the next round after updating the parameters to w'^T , we want to get closer to the correct answer (**be greater**):

$$w^T x < w'^T x$$

- Verify if it would do better after updating w as $w' = w + x$.

$$\begin{aligned} w'^T x &= (w + x)^T x \\ &= w^T x + x^T x \\ &= w^T x + \|x\|^2 \end{aligned}$$

- Since $\|x\|^2 > 0$, the condition is verified.

Perceptron: Learning Algorithm

- **2nd case:** $x \in C_2$, but was classified in C_1 . In other words, the correct answer is 0, which corresponds to $w^T x < 0$, but the model provides $w^T x \geq 0$. In the next round after updating the parameters to w'^T , we want to get closer to the correct answer (**be smaller**):

$$w^T x > w'^T x$$

- Verify if it would do better after updating w as $w' = w - x$.

$$\begin{aligned} w'^T x &= (w - x)^T x \\ &= w^T x - x^T x \\ &= w^T x - \|x\|^2 \end{aligned}$$

- Since $\|x\|^2 > 0$, the condition is verified.

The Perceptron Algorithm: Example


Example: $(-1, 2) -$ ✗
 $(1, 0) +$ ✓
 $(1, 1) +$ ✗
 $(-1, 0) -$ ✓
 $(-1, -2) -$ ✗
 $(1, -1) +$ ✓

$$w_1 = (0, 0)$$

$$w_2 = w_1 - (-1, 2) = (1, -2)$$


$$w_3 = w_2 + (1, 1) = (2, -1)$$

$$w_4 = w_3 - (-1, -2) = (3, 1)$$



Representational Power of Perceptrons


- Marvin Minsky and Seymour Papert, "Perceptrons" 1969: **The perceptron can only solve problems with linearly separable classes.**
 - The functions can be drawn in 2-dim graph and a single straight line separates values in two parts.
- Examples of linearly separable Boolean functions:



Representational Power of Perceptrons

- The perceptron cannot model the non-linearly separable functions:
Logical XOR function (computes the logical exclusive).
 - Input: Two input arguments with values in $\{0,1\}$.
 - Output: 1 if and only if two inputs have different values.
 - For such functions, we have to use multi-layer feed-forward network.

X_1	X_2	Y
0	0	0
0	1	1
1	0	1
1	1	0



Summary of The Perceptron Algorithm

- For a random sequence x_1, x_2, \dots, x_k , with $x_i (i = 1, \dots, k) \in C_1$ or C_2
- For each x , if it is correctly classified, then $w_{k+1} = w_k$, otherwise,

$$w_{k+1} = \begin{cases} w_k + x & \text{if } x \in C_1 \\ w_k - x & \text{if } x \in C_2 \end{cases}$$

- Convergence theorem: regardless of the initial choice of weights, if the two classes are linearly separable, there exists w such that:

$$\begin{cases} w^T x \geq 0 & \text{if } x \in C_1 \\ w^T x < 0 & \text{if } x \in C_2 \end{cases}$$

- The learning rule will find such solution after a finite number of steps.

Model Selection

What is Model Selection?

Given a set of models $M = \{M_1, M_2, \dots, M_R\}$, choose the model that is expected to do the best on the **test data**. M may consist of:

- Same learning model with **different complexities** or **hyperparameters**.
 - Nonlinear regression: polynomials with different degrees
 - K-Nearest Neighbors: Different choices of K
 - Decision Trees: Different choices of the number of levels/leaves
 - SVM: Different choices of the misclassification penalty
 - Regularized models: Different choices of the regularization parameter
 - Kernel based methods: Different choices of kernels ...and almost any learning problem
- Different **learning models** (e.g. SVM, kNN, DT, etc)
- Note: usually considered in supervised learning but unsupervised learning faces this issue too.

Held-out Data

- Set aside a fraction (10-20%) of the training data.
- This part becomes our held-out data (validation/development)
- Remember: Held-out data is NOT the test data
- Train each model using the remaining training data
- Evaluate error on the held-out data
- Choose the model with the smallest held-out error
- Problems:
 - Wastes training data
 - If there was an unfortunate split (can be alleviated by repeated random subsampling)

Cross-Validation

- K -fold Cross-Validation on N training examples
- Create K equal sized partitions of the training data
- Each partition has N/K examples
- Train using $K - 1$ partitions, validate on the remaining partition
- Repeat the same K times, each with a different validation partition
- Choose the model with the smallest average validation error
- Usually K is chosen as 5 or 10.

Leave-One-Out (LOO) Cross-Validation

- Special case of K -fold Cross-Validation when $K = N$
- Each partition is now **an example**
- Train using $N - 1$ examples, validate on the remaining example
- Repeat the same N times, each with a different validation example
- Choose the model with the **smallest average validation error**
- **Can be expensive** for large N . Typically used when N is small

Random Subsampling Cross-Validation

- Randomly subsample a fixed fraction αN ($0 < \alpha < 1$) of examples; call it the validation set
- Training using the rest of the examples, measure error on the validation set
- Repeat K times, each with a different randomly chosen validation set
- Choose the model with the smallest average validation error
- Usually α is chose as 0.1, K as 10

Bootstrapping

- Given a set of N examples
- Idea: Sample N elements from this set with **replacement** (already sampled elements can be picked again)
- Use this new set as the training data
- The set of examples not selected as the validation data
- For large N , training data consists of about only **63% unique** examples

- **Expected model error:**

$$e = 0.632e_{test} + 0.368e_{training}$$

- This can break down if we overfit and $e_{training} = 0$

Information Criteria Based Methods

- Akaike Information Criteria (AIC)

$$AIC = 2k - 2\log(\mathcal{L})$$

- Bayesian Information Criteria (BIC)

$$BIC = k\log(N) - 2\log(\mathcal{L})$$

where

- k : # of model parameters
- N : # of data examples
- \mathcal{L} : maximum value of the model likelihood
- Applicable for probabilistic models
- AIC/BIC penalize model complexity; the smaller the better.

Feature Selection

Selecting a useful subset from all the features. Why?

- Some algorithms **scale (computationally) poorly** with increased dimension
- **Irrelevant** features can confuse some algorithms
- **Redundant** features adversely affect regularization
- Removal of features can **increase (relative) margin** (and generalization)
- Reduces data set and resulting model size
- Note: Feature Selection is different from Feature Extraction
 - The latter transforms original features to get a small set of new features (we will discuss in dimensionality reduction).

Feature Selection Methods

- Methods agnostic to the learning algorithm
 - Preprocessing based methods
 - E.g., remove a binary feature if it's ON in very few or most examples
 - Filter Feature Selection methods
 - Use some ranking criteria to rank features
 - Select the top-ranking features
- Wrapper Methods (keep the learning algorithm in the loop)
 - Requires repeated runs of the learning algorithm with different set of features
 - Can be computationally expensive

Filter Feature Selection

- Uses heuristics but is much faster than wrapper methods
- **Correlation Criteria**: Rank features in order of their correlation with the labels

$$R(X_d, y) = \frac{cov(X_d, y)}{\sqrt{var(X_d)var(y)}}$$

- **Mutual Information Criteria**

$$MI(X_d, y) = \sum_{X_d \in \{0,1\}} \sum_{y \in \{-1,1\}} P(X_d, y) \log \frac{P(X_d, y)}{P(X_d)P(y)}$$

- High mutual information means high relevance of that feature
- These probabilities can be easily estimated from the data

Wrapper Methods

- Forward Search
 - Start with no features
 - Greedily include the most relevant feature
 - Stop when selected the desired number of features
- Backward Search
 - Start with all features
 - Greedily remove the least relevant feature
 - Stop when selected the desired number of features
- Inclusion/Removal criteria uses cross-validation

Summary of Today's Lecture

- Generative vs Discriminative Classification
- Linear Discriminant Analysis
- The Perceptron Algorithm
- Model Selection